



Building Your ITS: From Concept to Code

A PRACTICAL GUIDE TO IMPLEMENTING YOUR FIRST INTELLIGENT TUTORING SYSTEM

CIS*3750 – SYSTEM ANALYSIS AND DESIGN IN APPLICATIONS

PRESENTED BY: RICARDO JARRIN

TUESDAY, OCTOBER 21, 2025

Agenda

- ▶ Part 1: The Foundation
 - ▶ Going over the basics of a learning system
 - ▶ Simulate system using rules-based flash card demo + code
- ▶ Part 2: The ML Menu
 - ▶ Going through the ML aspects presented in the Sept 16 Lecture
 - ▶ Simple programs meant to give you an idea of how to implement
- ▶ Part 3: Live Demo
 - ▶ Going through an ITS that uses knowledge tracing + code
- ▶ Notable Resources
- ▶ Q&A

Part 1: The Foundation

A RULES-BASED ITS

The Goal for Today

- ▶ You understand the **what** – the four main ITS components. Today, we focus on the **how**
- ▶ We'll build a complete, rules-based tutor from scratch, step-by-step
- ▶ Then, we'll show how to “level up” **any** of the four components with specific Machine Learning tools and code examples

The Simplest Analogy: A ‘Smart’ Flashcard Deck

- ▶ A regular deck is **static**. A “smart” deck (our ITS) pays attention and adapts
- ▶ It knows the topics (**Domain**), tracks your progress (**Student**), decides what to show you next (**Tutoring**), and handles the interaction (**UI**)

Our Case Study: A Simple Algebra Tutor

- ▶ **Domain:** Solving basic equations
- ▶ **Skills:** addition, multiplication, solving linear-equations
- ▶ **Dependency:** addition → multiplication → solving linear-equations
- ▶ We will now build this tutor using a simple, non-ML approach

Walkthrough Step 1: The Domain and Student Models (The Data)

- ▶ The Rulebook (Domain): We'll use simple JSON files to store out skills, questions, and dependencies.
- ▶ The Scorecard (Student): We'll use a separate JSON file for each student to track their mastery score

```
{  
  "addition": [{"q": "5 + 8", "a": "13"}],  
  "multiplication": [{"q": "3 * 7", "a": "21"}],  
  "solving-linear-equations": [{"q": "2x + 3 = 7", "a": "2"}]  
}
```

```
{  
  "multiplication": ["addition"],  
  "solving-linear-equations": ["multiplication"]  
}
```

```
{  
  "student_id": "alex",  
  "mastery": {  
    "addition": 0.0,  
    "multiplication": 0.0,  
    "solving-linear-equations": 0.0  
  }  
}
```

Walkthrough Step 2: The Tutoring & UI Models (The Logic)

- ▶ The Game Master (Tutoring): The core of our system will be a Python function with if/else logic
- ▶ The Game Console (UI): We'll start a Command-Line interface (CLI) using Python's print() and input() functions.

Walkthrough Step 3: The Code - Tying It All Together

The Game Master →
tutor.py

```
tutor.py
1  MASTERY_THRESHOLD = 0.8 # Define what "mastered" means
2
3  def choose_next_skill(student_model, domain, dependencies):
4      """Finds the next skill to teach."""
5      mastery = student_model.get('mastery', {})
6
7      # Iterate through all available skills
8      for skill in domain.keys():
9          # 1. Check if the skill is already mastered
10         if mastery.get(skill, 0.0) >= MASTERY_THRESHOLD:
11             continue
12
13         # 2. Check if prerequisites are met
14         prereqs = dependencies.get(skill, [])
15         prereqs_met = all(mastery.get(p, 0.0) >= MASTERY_THRESHOLD for p in
16                           prereqs)
17
18         # 3. If not mastered AND prereqs are met, this is our skill!
19         if prereqs_met:
20             return skill
21
22
23  return None # No unmastered skills with met prerequisites
```

The Main Loop → main.py

```
main.py
1  def run_tutor(student_id):
2      """The main engine of the ITS."""
3
4      # 1. Load all data models from the 'data/' folder
5      domain, deps = load_domain('data/')
6      student = load_student_model('data/', student_id)
7
8      # 2. Start the main tutoring loop
9      while True:
10          # 3. TUTOR MODEL makes a decision
11          next_skill = choose_next_skill(student, domain, deps)
12
13          if next_skill is None:
14              print("Congratulations! You've mastered everything!")
15              break
16
17          # 4. UI MODEL presents a problem
18          problem = random.choice(domain[next_skill])
19          answer = input(f"Question on {next_skill}: {problem['q']} -> ")
20
21          # 5. Update STUDENT MODEL and save progress
22          is_correct = (answer.strip() == problem['a'])
23          student = update_mastery(student, next_skill, is_correct)
24          save_student_model('data/', student)
```



Part 2: The “ML Menu”

UPGRADING YOUR TUTOR

Level Up: Where to Add Machine Learning

- ▶ Our rules-based tutor works, but ML can make it smarter, more automated, and more adaptive
- ▶ Now, we'll spotlight one ML upgrade for each of the four components, complete with runnable code examples

ML for the Domain Model: Skill Discovery with *scikit-learn*

- ▶ **Problem:** Manually tagging hundreds of questions with the correct skill is tedious and doesn't scale well.
- ▶ **ML Solution:** Use unsupervised clustering to automatically group similar questions, helping to discover the underlying skills in your content.
- ▶ **Tool:** scikit-learn, a fundamental Python library for machine learning.

ML for the Domain Model: Difficulty Estimation

- ▶ **Problem:** We don't know if a new question is too easy or too hard.
- ▶ **ML Solution:** Train a regression model to predict a question's difficulty based on its features (e.g., length, word complexity).
- ▶ **Tool:** scikit-learn.

ML for the Student Model: Knowledge Tracing with pyBKT

- ▶ **Problem:** Our simple +/- score for student mastery isn't very nuanced. It doesn't account for lucky guesses or accidental slips
- ▶ **ML Solution:** Use Bayesian Knowledge Tracing (BKT) to model student knowledge as a probability, providing a more realistic estimate of their mastery. This is a form of knowledge tracing
- ▶ **Tool:** pyBKT, a Python library dedicated to BKT modeling

ML for the Student Model: Recommendation

- ▶ **Problem:** Our tutor's rule for "what's next" is basic. It doesn't know what *similar* students found helpful.
- ▶ **ML Solution:** Use collaborative filtering to recommend the next skill based on the success patterns of other students.
- ▶ **Tool:** scikit-surprise

ML for the Tutoring Model: Feedback Personalization

- ▶ **Problem:** Our rules-based tutor's feedback is generic. It can't adapt its teaching style to what works best for an individual student
- ▶ **ML Solution:** Use a classification model to predict which type of feedback (e.g., a text hint vs. a worked example) is most likely to help a student, given their current state. This is a form of feedback personalization
- ▶ **Tool:** scikit-learn, for training a simple and effective classification model

ML for the Tutoring Model: Policy Learning

- ▶ **Problem:** Our "if/else" logic is just a simple strategy. What is the *optimal* teaching policy to maximize long-term learning?
- ▶ **ML Solution:** Model the tutor as an "agent" in an environment. Use Reinforcement Learning (RL) to let the agent learn the best policy over thousands of simulated sessions.
- ▶ **Tool:** gymnasium

ML for the Tutoring Model: Dialogue Systems

- ▶ **Problem:** Our tutor can't have a conversation. It can only give quiz questions
- ▶ **ML Solution:** Use a pre-trained conversational LLM to power a chatbot that can answer student questions in natural language.
- ▶ **Tool:** transformers

ML for the UI Model: Affective Computing with *TextBlob*

- ▶ **Problem:** The system is unaware of the student's emotional state. A frustrated student may be about to quit
- ▶ **ML Solution:** Use Natural Language Processing (NLP) to perform sentiment analysis on student input to detect frustration or other emotions. This is a form of affective computing
- ▶ **Tool:** *TextBlob*, a user-friendly Python library for common NLP tasks

Reminder of ML Aspects Covered Before

- ▶ Domain
 - ▶ Automatic Skill Discovery
 - ▶ Difficulty Estimation
 - ▶ Knowledge Tracing
- ▶ Student
 - ▶ Prediction
 - ▶ Classification
 - ▶ Recommendation
 - ▶ Clustering
- ▶ Tutoring
 - ▶ Policy Learning
 - ▶ Dialogue Systems
 - ▶ Feedback Personalization
 - ▶ Adaptive Hinting
- ▶ UI
 - ▶ Adaptive Feedback Presentation
 - ▶ Affective Computing
 - ▶ Personalization



Part 3: Live Demo

THE STUDENT JOURNEY

Additional Tools & Resources

- ▶ Domain Model (The Rulebook)
 - ▶ Neo4j: A graph database for modeling complex skill dependencies.
 - ▶ RDFLib: A Python library for creating formal knowledge representations (ontologies).
- ▶ Student Model (The Scorecard)
 - ▶ SQLite: A simple, file-based database (built into Python) to manage multiple students.
 - ▶ pyKT: A library for more advanced Deep Learning Knowledge Tracing models.
- ▶ Tutoring Model (The Game Master)
 - ▶ Rasa: An open-source framework for building conversational AI and chatbots.
 - ▶ Gymnasium (formerly OpenAI Gym): A toolkit for developing Reinforcement Learning policies.
- ▶ UI Model (The Game Console)
 - ▶ Plotly Dash: A Python framework for building interactive dashboards and data visualizations.
 - ▶ Django: A full-featured Python web framework for building large, robust applications.
- ▶ Project Management
 - ▶ Git & GitHub: The industry standard for version control and team collaboration.

Public Datasets for Your Project

- ▶ Mathematics
 - ▶ MATH Dataset: Challenging competition math problems with step-by-step LaTeX solutions.
 - ▶ GSM8K: Grade-school math word problems with reasoning and final answers.
- ▶ Science & General Knowledge
 - ▶ SciQ: Science exam questions with multiple-choice answers and support paragraphs that serve as explanations.
 - ▶ OpenBookQA: Science questions based on a set of open-book facts.
- ▶ Programming
 - ▶ Project CodeNet: A massive dataset of programming problems and millions of user-submitted solutions.
 - ▶ APPS: A dataset of programming problems with test cases for validation.
- ▶ Pro-Tip for Your Project:
 - ▶ For the scope of a class project, creating your own small, high-quality dataset is often better than trying to parse a massive public one. This gives you full control over the content and is a key part of modeling your domain.

Your Project Roadmap & Final Advice

► The Foundation

- ▶ Pick a **very small domain** (e.g., 3-4 related concepts).
- ▶ Build your Domain Model using **JSON files**.
- ▶ Create a simple **Command-Line Interface (CLI)** for your UI.

► The Logic

- ▶ Implement a rules-based Tutoring Model with if/else logic.
- ▶ Build the basic Student Model to track scores in a JSON file

► Connect & Test

- ▶ Integrate all four components.
- ▶ Your goal is to have a complete, working, rules-based tutor.

► After: Upgrade with ML

- ▶ Pick one component and upgrade it using one of the ML techniques we discussed.



Questions?